Carbon Pricing and Trade Diversion

Jannik Hensel* Giacomo Mangiante† Luca Moretti‡

October 2023

Abstract

This paper examines the impact of carbon pricing on the trading patterns of firms and the overall effect on CO_2 emissions embedded in imports. We exploit carbon policy shocks, which are identified using high-frequency identification in combination with French administrative trade data and CO_2 emission databases. We show that a tighter carbon pricing regime of the European Union Emission Trading System disproportionately increases imports from countries outside of the carbon price domain. However, this effect is not permanent and does not lead to persistent changes in trading patterns. We document that while CO_2 emissions embedded in imports increases, this increase is slower than the response in value of trade, and firms mainly subsitute towards less CO_2-intensive inputs. Finally, we show that policies for specific, mostly carbon-intensive, products and industries aimed at preventing carbon leakage are successful in their objective.

Keywords: Climate policies, Carbon pricing, International Trade.
JEL classification: F14, F18, Q43, Q54

---

*University of Zurich, Zurich, Switzerland. Email: jannik.hensel@econ.uzh.ch
†Bank of Italy, Rome, Italy. Email: mangiante.giacomo@gmail.com
‡University of Zurich, Zurich, Switzerland. Email: luca.moretti@econ.uzh.ch
We would like to thank David Dorn, Nir Jaimovic, Florian Scheuer, and Josef Zweimüller for their valuable comments and guidance. We are particularly indebted to Diego R. Känzig for providing the data on carbon policy surprises. We are also grateful to Florin Bilbiie, David Hémous, Ralph Ossa, Jean-Paul Renne, and seminar and conference participants at the UZH Macro Lunch for helpful comments. Funding for this project was provided by UZH's University Research Priority Programs Equality of Opportunity. Access to some confidential data, on which this work is based, has been made possible within a secure environment offered by CASD – Centre d'accès sécurisé aux données (Ref. 10.34724/CASD). The views expressed in this paper are those of the authors and do not necessarily reflect those of Banca d'Italia or the Eurosystem.

1 Introduction

Limiting climate change and cutting greenhouse gas emissions is a significant challenge for humanity in the 21st century. While a general rise in average temperature seems unavoidable, the actual increase is still to be determined and ultimately depends on today's policies. Economists consider carbon pricing as one of the most, if not the most, efficient tools in limiting greenhouse gas emissions following a classical Coasian argument. Such a price can be implemented either with a direct tax on carbon emissions or through a cap-and-trade system. Such policies would ensure an efficient path of emmisions reductions. A global price on carbon emissions would not only limit the overall emissions but would ensure that the emissions with the lowest abatement cost are cut.

However, global policy coordination has proven to be notoriously complicated. Therefore, as a first step, several unilateral approaches have emerged, notably the European Union Emission Trading System (EU-ETS). However, they are usually incomplete since they lack a Carbon Border Adjustment Mechanism, and thus their effectiveness is potentially threatened in a globalized world where production is mobile. Already, a substantial part of CO_2-emissions is embedded in trade, and according to OECD statistics, gross imports of the EU-27 in 2018 contained a total of 1.2 billion tonnes of CO_2. In the same period, total production in the EU itself emitted about 3.4 billion tonnes of CO_2. Hence a substantial part of the latter emission is priced under the EU-ETS, this does not apply to the imports. Intuitively, this makes production within the EU-ETS more expensive, and the relative price of imports decreases, making them more attractive. This overall mechanism of CO_2 price evasion and its ecological consequences is often called carbon leakage. However, this is not conclusive evidence and while the theoretical mechanism sounds straightforward there is an ongoing academic debate abouts its practical relevance .

This paper aims to contribute to this discussion, by analyzing how changes in the tightness of carbon policies and resulting carbon price reactions of the EU-ETS affect trading patterns. For this, we assess both the effects on the value of imports as well as the overall CO_2 emissions embedded in these trade flows. We measure exogenous changes in the carbon price using the carbon policy shock series developed in Känzig (2022).

To evaluate how firms' trade flows are affected by carbon pricing policies, we combine the carbon policy shock series with French firm-level administrative and customs data. We extend the analysis to overall emissions by using both Exiobase and Ecoinvent data. Exiobase and Ecoinvent are datasets that aim to capture the CO_2 emissions embedded in production on a product-country level. While Exiobase is more aggregate and exhaustive, Ecoinvent is more granular and focuses on intermediate products. They complement each other in many dimensions which we outline later in the paper. The empirical specification we adopt is a panel local projection à la Jordà (2005), which allows to estimate fimrs' impulse response functions and trace the exact timing of their actions.

We document that firms' import decisions respond to carbon policy shocks. Firms import substantially more from countries not part of the EU-ETS following a carbon policy surprise. In contrast, trade from EU-ETS member countries stays flat or even declines. However, these effects are temporary and fade out. Quantitatively, this means that the firms' imports from countries not participating in the EU-ETS increase by up to 3.5% about three years after a shock that increases energy prices by 1% due to an increase in carbon prices. To further investigate carbon leakage, we show that the response in trade volume is also mirrored by CO2 emissions embedded in these import flows. Interestingly, this increase is weaker, suggesting that less CO_2-intense products drive the trade response—those emissions increase by up to 2.5% about three years after the shock. We analyze a set of policy measures that are in place for a subset of carbon-intensive industries/products with the aim of limiting carbon leakage. We confirm that these measures indeed help to attenuate import substitution and, through that, contribute to the muted response of CO_2 emissions. However, these policies come at a cost, e.g., free handouts of emission certificates. Therefore, despite these measures fulfilling their purpose, our results highlight the need to implement policies such as a

Carbon Border Adjustment Mechanism to limit trade reactions and to increase the overall economic as well as environmental efficiency of the EU-ETS.

Related literature Our first contribution to the literature is to estimate carbon leakage in response to environmental regulation. Carbon leakage is a particular case of the pollution haven hypothesis. Although there is an extensive literature on the pollution haven literature, the empirical evidence for carbon leakage is still limited, and the existing evidence of carbon leakage is mixed, depending on the context and specific study. Furthermore, there are different ways in which carbon leakage is defined in the literature, making comparisons more difficult. Typically, carbon leakage is measured as additional emissions abroad relative to domestic emissions savings. Carbon leakage of 100% means that all the emissions saved domestically are instead emitted abroad, which brings the global savings in emissions to zero. Carbon leakage can happen through different channels, e.g., through emissions embedded in trade but also through shifting energy demand affecting global energy prices (Yu et al., 2021). In this paper, we focus on the trade channel. Previous studies have used different techniques and data and also did not estimate carbon leakage as the percentage increase in embedded trade emissions using this level of granularity.

There is previous work that estimates carbon leakage under the EU-ETS. Muûls et al. (2022) compare firms that, due to their industry and installations, directly fall under the EU-ETS and consequently have to buy emission certificates with unregulated firms using a matched difference-in-difference design. They find no evidence of carbon leakage but a substantial reduction in CO_2 intensity. They argue that this can be rationalized in a framework of firms inattentive to the returns on investment in clean technology and, therefore, underinvest.

A key difference of our strategy is to use local projections and aggregate carbon policy surprises to the difference-in-difference strategy in this paper is that our approaches captures general equilibrium effects. Further, we argue that especially for imports the difference is not to expected between companies that fall under the EU-ETS and those who don't. But rather are product that are prdocued by firms that fall under the EU-ETS imported more.

Several studies investigate carbon leakage of emission-intensive, trade-exposed industries in the EU-ETS and generally find small or zero carbon leakage (see, e.g., Venmans et al., 2020). Dechezleprêtre et al. (2022) investigate the effect on multinational firms, where we would expect that emissions are more easily moved to their plants abroad. They find no evidence for a reallocation of CO_2 emissions.

Other studies analysed the EU-ETS focusing on different mechanism and channels of Carbon Leakage. **?** analyze the effect of the EU ETS on carbon leakage, but focusing on cross-country investments. They show that investments react to carbon pricing and that the effect is stronger for more polluting investments. However, the aggregate amount of diverted investments is small. Our paper in contrast looks at the trade channel of Carbon Leakage. In their review of the literature Venmans et al. (2020) find robust positive effects on innovation, especially on patenting activity.
But the literature is not restricted to the EU-ETS. For example, by running country-sector regressions Aichele and Felbermayr (2015) estimate carbon leakage in reaction to the overall policy packages implemented by countries participating in the Kyoto Protocol ratification. The Kyoto Protocol was implemented with different policies across countries. The authors estimate that imports from non-ratifying countries increase by 5%, which results in an 8% higher CO_2 content. While this paper therefore assess the effect of introducing a policy package consisting of a wide range of measures, our approach allows to estime the effect of marginally adjusting an existing policy.

Another approach to estimating carbon leakage is to use computational general equilibrium models (CGE). This approach has the advantage of potentially taking several channels of carbon leakage into account. In a review Yu et al. (2021) find that typical carbon leakage estimates in this literature are between 5 and 30%, with both the energy price channel and the trade channel playing an important role. This models have the advantage of explicitly modeling the introduciong and potential channels of carbon prices. Delivering an exogenous measure for firm's trade reaction to carbon policy shocks can be used as an input to improve the fit of these models.

Känzig (2022) developed the shocks used in this paper. He shows that low-income households are over-proportionally affected by carbon price shocks, mainly through general equilibrium effects, and also documents more green innovation. Building on this work, several paper have already emerged extending this analysis. Mangiante (2023) document that poorer Euro Area countries' economic activity is more strongly affected. Berthold et al. (2023) show that economic activity in more CO_2 intensive countries is more affected by carbon pricing policies. Hensel et al. (2023a) outline that firms' inflation expectations increase significantly in reaction to the shock. However, firms overestimate the effect, which leads to a positive forecast error in the medium term.

In contrast to previous empirical studies on carbon leakage that rely, e.g., on a difference-in-difference framework or similar identification methods, our empirical approach captures the general equilibrium effects of carbon price changes. Primarily through the spillover effects of energy prices, they not only affect companies who are directly responsible for emissions and are subject to the emission trading system but essentially all companies through their energy consumption. Further, using linear projections allows us to trace firm-level carbon leakage responses over time.

Road map. The paper is structured as follows. Section 2 provides a comprehensive overview of the relevant features of the EU-ETS, chapter 3 outlines the data used, Chapter 4 shows the effects of the Carbon Price Shocks on key macroeconomic variables, Chapter 5 presents the empirical strategy, and Chapter 6 summarises the main results. Chapter 7 discusses potential caveats, and chapter 8 concludes.

2 The EU-ETS

> *"The EU-ETS is a cornerstone of the EU's policy to combat climate change and its key tool for reducing greenhouse gas emissions cost-effectively."*
>
> EU Comission

The EU-ETS operates under a cap and trade system, meaning that not the price but the quantity of greenhouse gas emissions is fixed. The price is then set through an auction mechanism. Companies that are regulated under the EU-ETS are required to obtain sufficient certificates for their greenhouse emissions. These are checked annually and are subject to fines if the account is not balanced. As long as the maximum quantity of emissions certificates is a binding constraint, this mechanism ensures a positive price. The system was introduced in 2005 and has undergone both significant changes and minor revisions over time. The major changes can be divided into 4 phases of the EU-ETS, with the first phase spanning from 2005-2007, phase 2 from 2008-2012, phase 3 from 2012-2020, and phase 4 from 2021-2030. The first phase was designed as a learning-by-doing stage, and most emission certificates were allocated for free. It was restricted to activities related to power generators and energy-intensive industries. The aim was to put the infrastructure and monitoring in place. Furthermore, there was not one cap for the overall ETS domain but for each country, which was based on estimates. These turned out to be too generous, and the price eventually fell to zero as they were also not allowed to be carried over to the second stage. A primary objective of this stage was to get data on total emissions, which could then be used to determine the cap in the subsequent phases. The second phase expanded the scope of activities, installations, and companies covered under the EU-ETS, reduced the amount of emission certificates, and reduced the share of allowances allocated for free. Phase 3 replaced the national caps with an EU-ETS-wide cap, replaced the free allocation of allowances for many industries with an auction mechanism, and further expanded the scope of covered activities. The EU Commission formally adopted the latest major reform of the EU-ETS in April 2023. The main target is the gradual phase-out of free allowances still granted to industries/products deemed to be at high risk of Carbon leakage. In parallel with the phase-out of free allowances, a so-called carbon border adjustment mechanism (CBAM) will be implemented to reduce the risk of carbon leakage. Under the CBAM, emissions embedded in imports will be taxed; however, exports out of the EU will not be subsidized. Hence, weighing the internalization of the external costs of carbon higher than the competitive disadvantage that European firms face on global markets. In addition, the reform consists of a more ambitious reduction path that aligns with the EU's new climate goals, the gradual inclusion of additional sectors, and the creation of separate new ETS for buildings, roads, transport, and fuels.

Today, about 45% of European greenhouse emissions are covered by the EU-ETS. The EU-ETS covers approximately 10,000 companies and installations. Globally, the EU-ETS only plays a minor role, which is decreasing over time, covering 2.81% of global emissions in 2019, down from 4.85% in 2005. This decrease in coverage is primarily due to the EU emitting a smaller share of global emissions in 2019.

Concerns for Carbon Leakage have been a central issue considered in the design of the EU-ETS. The European Commission incorporated measures in the cap and trade system intending to prevent carbon leakage. As a result, specific products/industries at risk of carbon leakage are identified. The EU Commission considers two characteristics to be the main determinants of carbon leakage risk: First, the additional costs incurred due to the implementation of the EU-ETS, and second, the sector's reliance on trade and exposure to international competition. The additional costs are measured relative to gross value added, considering if an industry would be eligible for a share of free allowances independent of its carbon leakage risk. Costs include direct costs of the emission allowances and indirect costs through, e.g., higher electricity prices. Reliance on trade and competitiveness is measured as the share of imports and exports of a sector relative to domestic production plus imports (Graichen, V., et al. 2017). Sectors (4-digit nace) or subsectors (6-digit CPA) are deemed at risk of carbon leakage if they satisfy one of the following criteria. [1]:

- If the share of additional costs exceeds 30%
- If the trade intensity exceeds 30%
- If the share of additional cost exceeds 5% *and* trade intensity is above 10%

Sectors and subsectors that satisfy one of these criteria are eligible to receive up to 100% of their required emission allowances for free. The exact share is determined using a formula incorporating production quantity and a benchmark value for that particular product (measured in emissions per tonne). It is worth noting that other industries were also eligible for free allowances, although for a lower share, which has been reduced annually since 2013.

[1]For a few sectors and subsectors a qualitative assessment was carried out.

In addition, countries are allowed to compensate specifc industries that face high indirect costs up to a certain threshold.

2.1 Carbon Policy Shock Series

We exploit the carbon policy surprise developed by Känzig (2022). The identification strategy and the mechanics of this shock are similar to the identification of monetary policy surprises. (see, among others, **?** and Nakamura and Steinsson, 2018). It relies on the availability of high-frequency price data around special events. Monetary policy shocks rely on asset price (e.g., bonds) fluctuations in a narrow time window around the announcements of monetary policy decisions. In case of the carbon policy surprises, the counterparts for monetary policy decisions are policy decisions by the European Parliament, the European Commission, or European Courts that affect the supply of carbon emission certificates. These include, e.g., decisions regarding the free allocation of certificates, the auction of allowances, and the use of international credits. The counterpart of asses price movements are changes in the future prices of carbon emission allowances. Känzig (2022) identifies 126 of those events from 2005 to 2019. These events reveal information to market participants, and when looking at the market reactions in a tight time frame around the event, these can be interpreted as an unexpected and exogenous price change solely driven by the event.

Surprises in carbon policy are calculated by examining fluctuations in the future prices of EU emission allowances (EUAs) within the ICE market, which is the most liquid future market. More specifically, these surprises are characterized by the variations in the EUR value of carbon prices in relation to the wholesale electricity price prevailing on the day preceding the specific event. The daily surprises are then summed to a monthly series. In months with no regulatory events, the series has a value of zero.

Since it is hard to argue that the identified events are exhaustive and cover all such relevant events, the resulting surprise series is prone to measurement error. Therefore, the standard in the literature is to use the surprise series as an IV for a Vector Auto Regression. This VAR covers the period from 1999 to 2019. It includes the energy component of the HICP (for which the Instrument is used), total GHG emissions, the

Table 1: Descriptive Statistics of carbon price surprises and the resulting shock

	N	Mean	Std. Dev.	Min	Max
Surprise	246	-0.023	0.221	-2.007	1.532
Shock	246	0.000	0.630	-1.333	1.335

Notes: Surprises are identified directly around the announcement windows, while shocks are the structural shock series identified through an external instrument VAR approach with the surprises as an instrument. Both were supplied by Känzig (2022)

headline HICP, industrial production, the unemployment rate, the policy rate, a stock market index, and the real effective exchange rate.

The carbon policy shocks are then extracted from the residuals of the monthly VAR (see Stock and Watson, 2018) and normalized to increase the energy component of the HICP by one percent on impact. While this may sound large at first sight, given the low level of Carbon prices, carbon policy surprises of this magnitude were observed in the time frame concerned. We refer to Appendix A.1 for further detail.

If the procedure is successful, this shock series is truly exogenous to everything else that affects the economy. While this is impossible to test and verify, Känzig (2022) conducts several checks. For one, they are mean zero and not serially correlated. They are also not forecastable by past macroeconomic or financial variables and do not correlate with other known structural shock measures, including oil, uncertainty, financial, fiscal, and monetary policy shocks.

Given that the EU-ETS only covers a small share of global GHG and was, over most of the period, the most extensive Trading System, it can be ruled out that the carbon policy surprised the price effects in other ETS systems.

3 Data

3.1 Firm Level Data

The central part of our analysis relies on French firm-level trade data. The French customs collects the data and makes it available at the firm - product - destination/origin level, aggregated to a monthly frequency. To obtain these data, we had to get approval

Table 2: Descriptive statistics of the trade data

	Median	Mean	Std. Dev.
Months in sample	57	86	79
Months of zero trade	20	41	50
Number of countries	3	6	9
Number of products	6	21	43
Value of trade (per obs)	11'102€	161'890€	3'876'868€

Notes: Summary statistics computed based on our selected sample. The only restriction is that a firm has to trade more than once.

from the Direction Générale des Douanes et des Droits Indirects (DGDDI). The data was then provided through CASD.

It contains information on 578,734 unique firms; however, we restrict the sample to firms that trade for at least two months - this leaves 387,704 distinct firms in the sample. As shown in table 2, the median firm trades with three countries and with six different products. The median firm has been in our sample for almost five years, of which, in one-third of the months, we do not observe any trade. The median monthly value of a firm's product-country trade is 11'000€ while the median is 161'890€. The fact that trade is skewed towards larger firms making larger transactions is a well-known fact. We use a crosswalk provided by Eurostat to get time-consistent CPA 2002 product codes to merge the data with Exiobase.

When working with data on the origin or destination of traded products, we are faced with a specific problem, namely, which country is actually reported as the origin country. For example, if a product is shipped from China to France, it will often enter the EU through Rotterdam or Antwerp. Depending on the data source, this French import from China will now show up as a French import from the Netherlands or Belgium, which would severely limit the usefulness of this data for our purposes. The French customs data on imports reports both countries, which allows us to differentiate cleanly between imports from other EU countries and imports from outside the EU. The same does not apply to the export data, where only the first and not the final destination is reported. For example, an export from France to China that goes through the port of Antwerp would be reported as an export to Belgium. As non-EU exports are an integral part of our analysis, this limits what we can do regarding the effect of carbon prices on French Exports. A detailed description of this issue and the data, in general, can be found in Bergounhon et al. (2018)[2].

We supplement the firm-level trade data with balance sheet information, social security data, and an energy consumption survey. The Energy survey is available for a subset of manufacturing firms and contains detailed information on their overall energy cost and usage, also divided by the energy source. The balance sheets are available for the universe of French firms for all relevant years (Ficus/Fare). We combine the energy survey and the balance sheets to compute efficiency measures, e.g., energy expenditures per value added. Balance sheet data is available only at a yearly frequency, which, in combination with our empirical design, reduces the variation in the carbon policy surprise that can be exploited considerably.

All these datasets can be linked through a firm identifier (Siren). However, this identifier does not coincide precisely with the definition of a firm that we have in mind, which relies on operational control. Especially in the mid-2010s, several large firms split their operations into different administrative units without actually splitting the firms in terms of control. To get a time-consistent firm definition, we follow the approach developed and used by De Ridder et al. (2022) and Burstein et al. (2020), which is based on previous work by Isabel Méjan see, for example, di Giovanni et al. (2014). [3].

3.2 Emissions Data

We use two different data sets to get data on the emission content of trade: Exiobase and Econinvent.

Exiobase 3 provides consistent environmentally extended multi-regional input-output tables for 44 countries and a "rest of the world aggregate." The construction of this database is described in Stadler et al. (2018) and has been widely used in economic research, e.g. by **?**. The data covers 200 product groups or 163 sectors. This data has many desired properties. One advantage of this data is the time series character of the data, covering each year from 1995 to 2019. Further, it contains a variety of environmental emissions such as water usage, energy usage, pollutants, and CO_2

[2]We would like to thank Isabel Méjean for making the cleaning code publicly available on her web page

[3]We thank them for kindly providing us with their code.

emissions per value of production. We will focus on the latter to estimate the carbon content of trade. According to this data, 64% of products-country observations have a higher CO_2 content than when produced in France. We apply a crosswalk provided by Exiobase to get to CPA 2002 codes that can be linked to the customs data. Due to its aggregate nature, one can link all the observations in the customs data to an aggregated product category in Exiobase.

Ecoinvent (v 3.1.9) is a life cycle inventory database. It covers a wide variety of products and contains detailed information on a multitude of pollutant emissions. Here again, we will focus on CO_2 emissions. It has a more specialized scope with regard to products and countries, but it provides us with a very high granularity and emissions per kg of a product. It has a particular focus on products used as intermediates. Currently, we only have access to the most recent version, which does not contain historical data. If a product appears in Ecoinvent, the information is on a very detailed level (six-digit CPA2015). We can match roughly one-third of our trade flows to Ecoinvent products, predominantly in manufacturing (75%) and some in the agricultural sector (10%).

These two sources are great complements. Exiobase has the advantage of more comprehensive coverage and yearly measures of CO_2 intensities. However, the measure might be too aggregated, and emission intensities are measured relative to the value of goods. Ecoinvent, on the other hand, provides us with the most recent estimates and covers only parts of our sample. However, emissions are measured relative to weight, and the products covered are measured with greater detail.

4 French Macroeconomic Variables and Carbon Policy Shocks

Because Känzig (2022) estimates are based on data for the EA-19 countries, we first want to check how the French economy reacts to these shocks. For this we estimate the following local projection à la Jordà (2005) at the national level:

$$y_{t+h} = \alpha_h + \beta_h CPShock_t + \sum_{p=1}^{P} \theta_h^p y_{t-p} + \epsilon_{t+h}, \tag{1}$$

for the quarters $h = 1, ..., 16$. y_{t+h} is the dependent variable at time $t + h$ and $CPShock_t$ are the carbon policy shocks at time t extracted from the Proxy-VAR. In the baseline specification, we include three lags of the dependent variable and we correct for autocorrelation using Newey and West (1987) standard errors.

The main dependent variables are the log of the Energy Consumer Price Index (CPI), the number of unemployed, the number of employed, and real GDP. The coefficient of interest is β_h which captures the response of the dependent variable to a carbon policy shock for each horizon h.

The responses of the dependent variables to a carbon price shock are reported in Figure 1. Following a carbon policy shock that results in a one percent increase of the HICP energy component on impact, all Macro variables react as expected. The shock indeed increases the Energy component CPI by roughly 1 percent, meaning that the response is consistent with Känzig (2022). We also see that real activity is affected. Unemployment takes two quarters to react but is then persistently higher for the remaining 4.5 years. Employment is the mirror image of unemployment, although slightly more noisy. Real GDP decreases after one and a half years and tends to be lower for the next two years. However, this effect is not significant.

While the Energy Component of the CPI only reacts for about five quarters before moving back to its original level, that does not mean that the price propagation of the shock through the overall economy is completed. Hensel et al. (2023b) analyze how the shock impacts inflation expectations and pricing decisions. Using survey data of manufacturing firms, they show that firms keep adjusting their prices up to 4 years after the shock.

Figure 1: The carbon policy surprise series

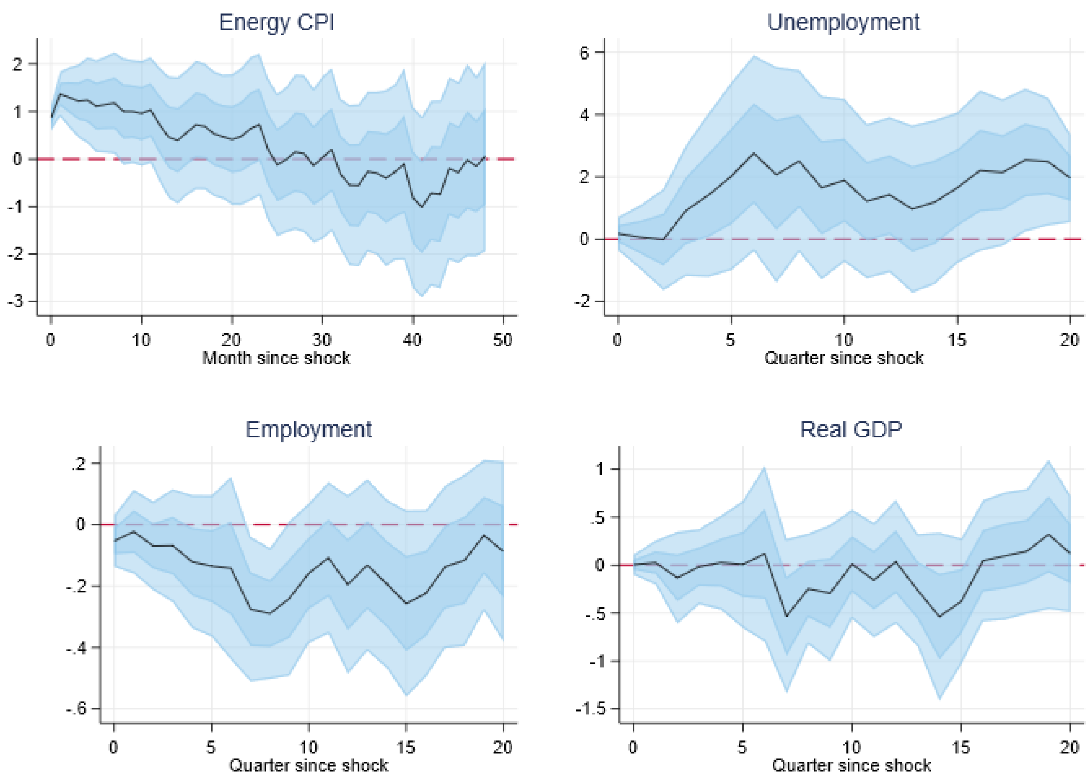

Notes: The figure plots the response to a carbon policy shock, normalized to increase the HICP energy by 1 percent on impact, for the French Energy CPI (top lef panel), the number of unemployed in France (top right panel), the number of employed people in France (bottom left panel) and the PPI (bottom right panel). The solid lines are the point estimate and the shaded areas are the 95 and 68 percent confidence bands, respectively.

5 Empirical Methodology

We have shown that aggregate measures of economic activity react following a carbon policy shock. We now shift our focus from macro- to firm-level variables.

We estimate the average firm-level response to a carbon policy shock following the approach used by Jordà (2005):

$$log(y^{i,c}_{t+h}) = \beta_h CPShock_t + \omega_h CPShock_t * ETS_c + \alpha^{i,c}_h + \sigma_{y,m} + \sum_{p=1}^{4} \theta^p_h X^i_{t-p} + \varepsilon^{i,c}_{t,h} \quad (2)$$

for horizon $h = 1, ..., 48$. $log(y_{t+h}^{i,c})$ is the dependent variable, e.g., six months ahead trade volume with country c, of firm i at time $t + h$ after the shock. ETS_c is a dummy equal to 1 if the respective country is a member of the EU-ETS and zero otherwise. β_h captures the effect of the carbon policy shock on trade. In contrast, ω_h captures how trade with a country that is also subject to the shock is affected differently. $\alpha_h^{i,c}$ are Firm-country-fixed-effects, $\sigma_{y,m}$ are year and month-fixed effects are X_{t-p}^i are time-varying firm characteristics (e.g. lagged imports). Finally, standard errors are two-way clustered at date and firm level. The coefficients of interest in this case are β_h and ω_h. In this specification, β_h can be interpreted as the reaction of non-ETS trade while the coefficient ω_h reports how ETS trade is affected differently. Therefore, the sum of β_h and ω_h reports the reaction of ETS-trade. The key to our identifying assumption is that the shocks, residuals to a proxy VAR, are uncorrelated to the error term. The shocks are normalized such that the coefficients reflect the reaction of trade to an increase in energy prices of 1% due to an increase in carbon prices.

As a first step, we analyze the response measured in terms of trade value. French administrative trade data reports the trade flows of each firm at the product/country level at a monthly frequency. The high frequency of the data and the long panel structure make it an ideal setting to study how firms' trading decisions are affected by changes in carbon price. One issue with this level of granularity is that most firms do not trade every month, and the frequency of zero trade flows in the original data. To reduce this kind of noise in our data, we define the log of the cumulative trade of the next six months as the dependent value of the value regressions. However, the frequency of our data and the impulse response function stays on a monthly level. While this increases the precision of our estimates, the point estimates without the six-month aggregations are qualitatively and quantitatively very similar.

In the second step, we aim to understand how the overall CO_2 emissions embedded in these trade flows react. For this purpose, we combine the product/country level trade data with the corresponding average CO_2 intensity from either Exiobase or Ecoinvent. We then aggregate them to obtain the overall CO_2 emissions embedded in all imports of a given firm in a given month. To avoid the problems with zeros mentioned in the prior paragraph, we apply the same procedures. We aggregate our measure of CO_2 embedded in trade over a six-month horizon and use this as a dependent variable in equation (2).

To further shed light on potential channels of heterogeneity, we divide the sample by sector or product classification. On a sectoral level, we distinguish between manufacturing and Wholesale and retail trade, repair of motor vehicles, motorcycles, and personal household goods. Both of these sectors are by far the most significant trading sectors and cover well above 90% of the trade volume. On a product level, we divide the sample by two-digit CPA2015 codes, e.g., Paper and Paper products, Chemicals and chemical products, or Basic pharmaceutical products and pharmaceutical preparations. As mentioned in section 2, several measures were implemented to limit carbon leakage in particularly exposed sectors. We investigate how these measures affected carbon leakage by including a dummy for these industries/industries and interacting it with the shock and the ETS dummy, including all possible combinations. We report the results from these regressions in section 6.2.

6 Results

6.1 Trade

We represent the baseline results of regression 2 with the value of trade as a dependent variable in Graph 2. The graphs can be interpreted as impulse response functions and report how trade reacts up to 48 months after the shock. The graph shows how trade reacts after an increase in energy prices of 1% due to an increase in carbon prices. The first graph reports the coefficients β_h in equation 2 and is the response in non-ETS trade. The results establish both a statistically and economically significant increase in imports in the medium term as a consequence of the shock. However, this increase is only temporary and fades out in the long run. Strikingly, this increase in non-ETS imports is not mirrored similarly by ETS country imports. The fact that non-ETS imports increase while ETS imports mainly stay flat is striking evidence for Carbon Leakage.

We can learn more about the nature of the shock and the adjustment process within firms by looking closer at the time profile of the response. In the first months after a

Figure 2: Effects on trade

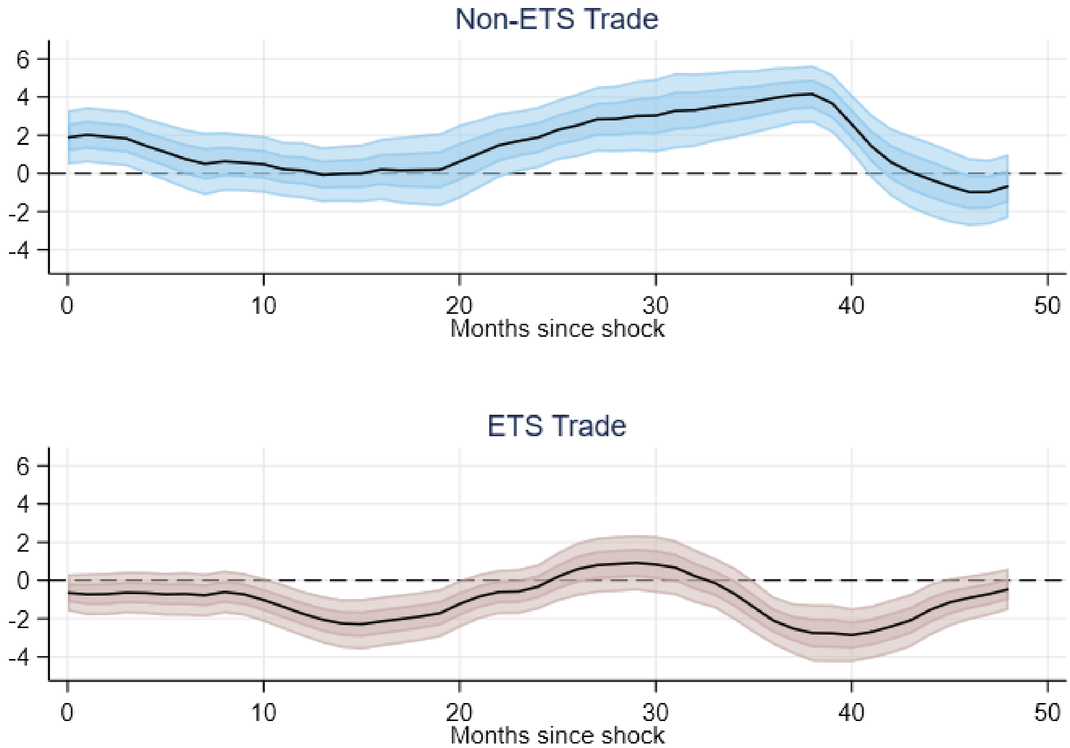

Notes: This figure displays the average firm response to a carbon price shock. The dependent variable is a 6-month moving average of the total value of trade. We include firm-country, year and month fixed effects and control for 4 lags of the dependent variable. Standard errors are clustered at the firm and date level. The solid lines are the point estimates, and the shaded areas are the 95 and 68 percent confidence bands, respectively.

carbon price surprise, there is very little movement in trade. This delay in the response of trade can have several reasons. For one, trade deals are often concluded ahead of time, and production and delivery times can add to these delays. Further sourcing new suppliers either with higher capacity or for new products takes time and usually involves some fixed costs. After about 18 months, firms started to import significantly more products from countries outside the EU-ETS domain. These results are in strong support for the carbon leakage hypothesis. However, over time, the effect of a carbon policy surprise on the value of trade reverts back to zero, which means that in the longer run, there is no sign of carbon leakage. We will show that these results are also accompanied by an increase in CO_2 embedded in imports in the next section.

The coefficients also highlight the economic significance of our findings. To put the results into perspective, non-ETS trade 30 months after the shock is about 3.5% higher than otherwise. This signals that companies in the medium run are very responsive to changes in CO_2 prices. The declining impact on trade is consistent with theories arguing that firms adjust to the higher carbon prices over time. For example, Muûls et al. (2022) argues that firms are inattentive concerning the return on investment in green technologies. Increasing carbon prices then incentivizes investments with positive returns. Another potential channel is that firms overestimate the persistence of the shock. Hensel et al. (2023a) show this is the case for price expectations. Känzig (2022) shows that patenting of low-carbon technologies increases after the shock, leading to greater availability of green technologies. The zero effect on trade, in the long run, is consistent with the previous findings in the literature that do not find permanent or only small carbon leakage in the EU-ETS.

We next turn to a subsample analysis and split the import data either by product categories or by sector. The sectoral response between manufacturing and retail, the two sectors responsible for about 90% of imports, is strikingly similar (put graph here). However, interesting patterns emerge on the product level. A sector that was a prime example in the literature of carbon leakage is paper and paper products (see, e.g., paper Aichele and Felbermayr, 2015), as these products also tend to be very carbon-intensive. Indeed, our findings support that carbon leakage in this product group tends to be stronger than in the baseline of all products. Furthermore, the response in ETS vs. non-ETS trade is more pronounced, too. Similar results can be found for chemicals and chemical products. Products where the response is notably muted or statistically insignificant are, for example, basic pharmaceutical products and pharmaceutical preparations. This pattern would suggest that product groups with higher CO_2 intensity show stronger reactions. However, overall, there seems to be no pattern that CO_2 intense products react more strongly, as we will point out in the next section.

6.2 Exemptions

In section 2, we described how carbon leakage was a primary concern when designing the EU-ETS and that measures were put in place to protect industries at high risk of carbon leakage. The core measure is to provide these 154, 4-digit industries with free emission allowances (up to 100%). In this section, we are investigating how these measures affect carbon leakage. To do so, we estimate the following regression at the firm level.

$$\begin{aligned} log(y_{t+h}^{i,c,p}) = \beta_h CPShock_t + \omega_h CPShock_t * ETS_c + \\ + \lambda_h EX_t * CPShock_t + \theta_h EX_t * CPShock_t * ETS_c + \\ + \gamma_t EX_t + \zeta_t EX_t * ETS_C + \alpha_h^{i,c,p} + \sigma_{y,m} + \sum_{p=1}^{4} \theta_h^p X_{t-p}^i + \varepsilon_{t,h}^{i,c} \quad (3) \end{aligned}$$

for horizon $h = 1, ..., 48$, country c, at time t of firm i for products that are deemed at risk of carbon leakage p. The dummy EX takes the value one if the product is deemed to be at risk of carbon leakage by the EU Commission. $log(y_{t+h}^{i,c})$ is the dependent variable, e.g., six months ahead of trade volume. β_h captures the effect the carbon policy shock has on trade. ω_h captures how trade with a country that is also subject to the shock is affected differently. λ_h captures how trade is affected differently if it is on the carbon leakage risk list, and ω_h captures how trade with ETS countries is affected

Figure 3: Trade patterns comparing exempt and non exempt products

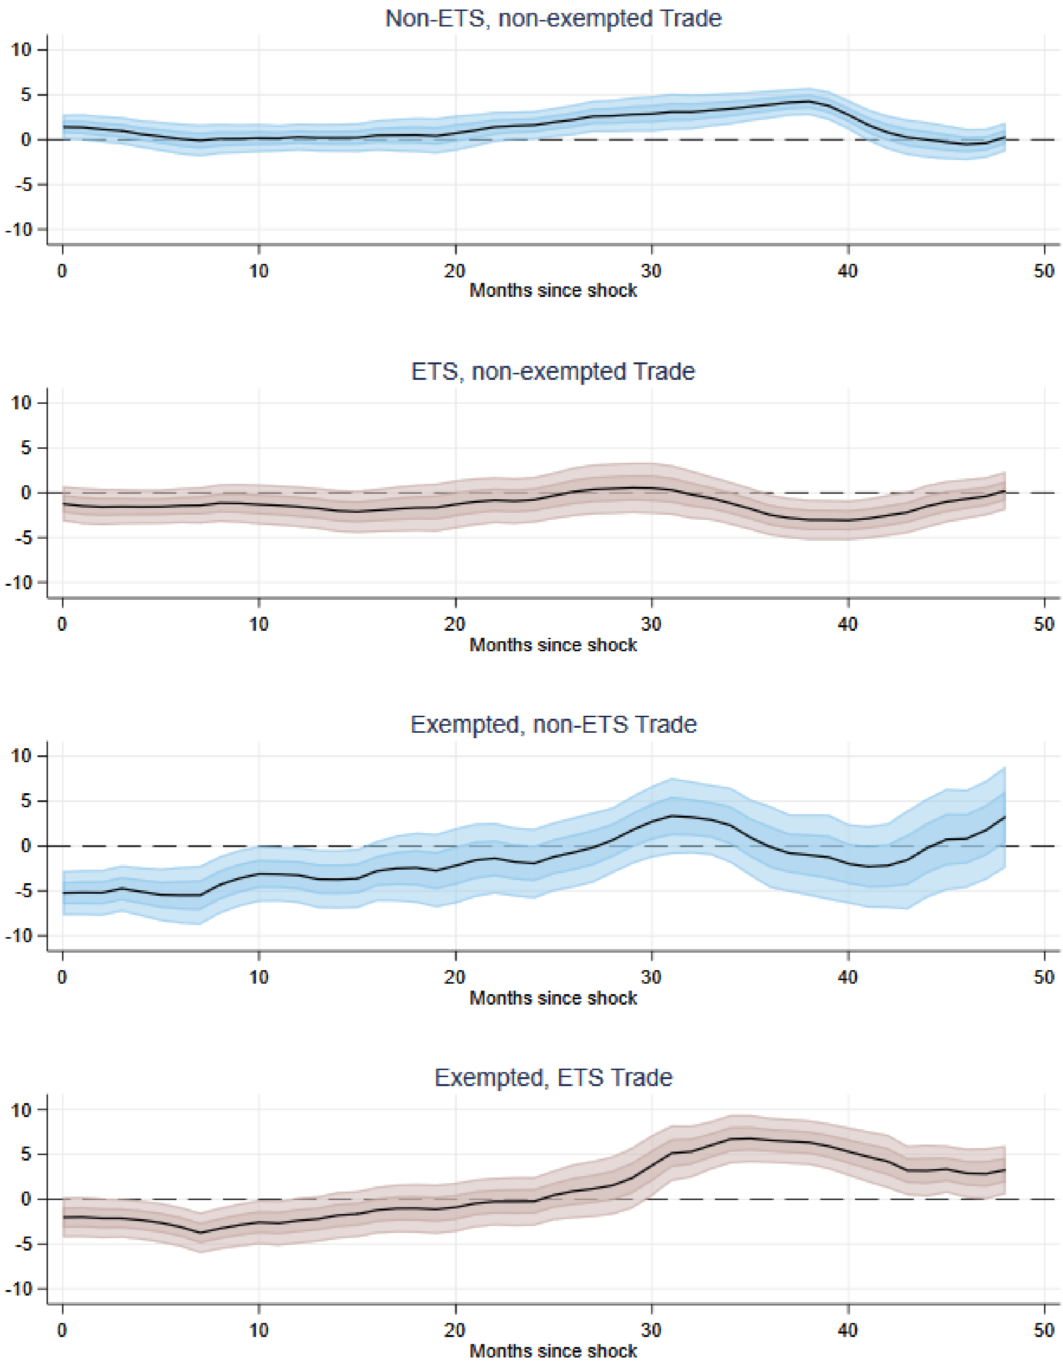

Notes: This figure displays the average firm response to a carbon price shock. The dependent variable is a 6-month moving average of the total value of trade. We include firm-country, year and month fixed effects and control for 4 lags of the dependent variable. Standard errors are clustered at the firm and date level. The solid lines are the point estimates, and the shaded areas are the 95 and 68 percent confidence bands, respectively.

differently if it is on the carbon leakage risk list [4]. For the reader's ease, we will report the absolute effects of the shock and leave the differential impacts in the appendix.

The rationale behind this specification is to identify those products that would potentially qualify for special treatment, such as free emission allowances, if they were produced in the EU-ETS domain and to check whether importing these goods from outside the EU-ETS behaves differently after the carbon price shock. While the allocation of free allowances does not change incentives, i.e., marginal costs of production, it decreases overall costs, increases overall profits, and therefore reduces the need to raise prices if international competition would make it hard to do so, at least in the short run. Most importantly, the latter mechanism is a specific reason for an industry or product to be included on the EU Commission's exemption list.

We depict the result of this regression with trade value as the dependent variable in Figure 3. Trade with non-exempted products reacts in the same way as aggregate trade. Trade within the EU-ETS is not significantly affected by the carbon price shock. Trade with non-ETS countries picks up after 18 months and remains elevated for the next two years, leading to significant carbon leakage during that time.

We show the impact of the carbon price shock on goods at significant risk of carbon leakage in the bottom two panels of Figure 3. Trade in these goods with non-ETS countries is less affected than other goods. We do not see the same increase after 18 months in non-exempt goods. The results are even more striking as trade competition is a key reason to be on the exemption list, and one would expect an even stronger result than for other goods.

The trade reaction with exempted goods within the EU-ETS shows an interesting pattern. We find that trade picks up similarly to trade in non-exempted goods with non-ETS countries, although the response is slightly delayed.

It seems that the measures are successful in protecting industries at risk from carbon leakage from competition from outside the EU-ETS. However, there is a particular reallocation within the EU-ETS away from French firms. As long as this reallocation is towards more efficient firms, this might be the ideal scenario for the European Union.

[4] As a reminder for the reader: $\alpha_h^{i,c}$ are Firm-country-fixed-effects, $\sigma_{y,m}$ are year and quarter-fixed effects are X_{t-p}^i are time-varying firm characteristics (e.g. lagged imports). Standard errors are two-way clustered on date and firm level.

6.3 Carbon Leakage

In this section, we focus on the embedded emissions, the dependent variable now being total emissions embedded in trade. We use information on emissions from two different databases, Exiobase and Ecoinvent. For this, we merge the product level data with the corresponding level in CO2 emission from these databases. Exiobase measures CO2 emissions embedded in value, so we multiply the trade value in our data with the CO2 measure provided to get the total CO2 content. Ecoinvent provides the CO2 content by kilogram, and therefore, we multiply the weight measure in our data with the CO2 measure provided by Ecoinvent. These measures are then aggregated on a firm and month level to get total CO2 emissions embedded in import for each firm and month.

In the top two panels of Figure 4, we use emissions data coming from Exiobase. The time profile is very closely aligned with the response of trade value. In the first periods, embedded emissions in trade do not react significantly. It is only over time that emissions from trade with non-ETS countries start increasing. The peak is reached after approximately three years when emissions embedded in trade with non-ETS countries are up to 2.6% higher than they would have been without the shock. After that, embedded emissions start decreasing again until they reach the level that would have been preserved without the shock.

In the bottom two panels of Figure 4, we plot the coefficients using emissions data from Ecoinvent. We have discussed the advantages and disadvantages of both emissions datasets extensively in section 3. It is comforting to know that our results are robust to the choice of emissions data. In particular, given that Exiobase is based on transaction value while Ecoinvent is based on transaction volume, which is, for instance, independent of any price or exchange rate movements. The time profile remains qualitatively the same, and the magnitudes coincide quantitatively. Any difference in timing might be due to the sample that changes between the datasets, with Ecoinvent having smaller coverage than Exiobase.

One important channel for the muted CO_2 response is the EU Commission's exemption list. Products on this exemption list are carbon intensive. In the previous section, we highlighted that the exemptions are successful in preventing carbon leakage, and

Figure 4: Carbon leakage

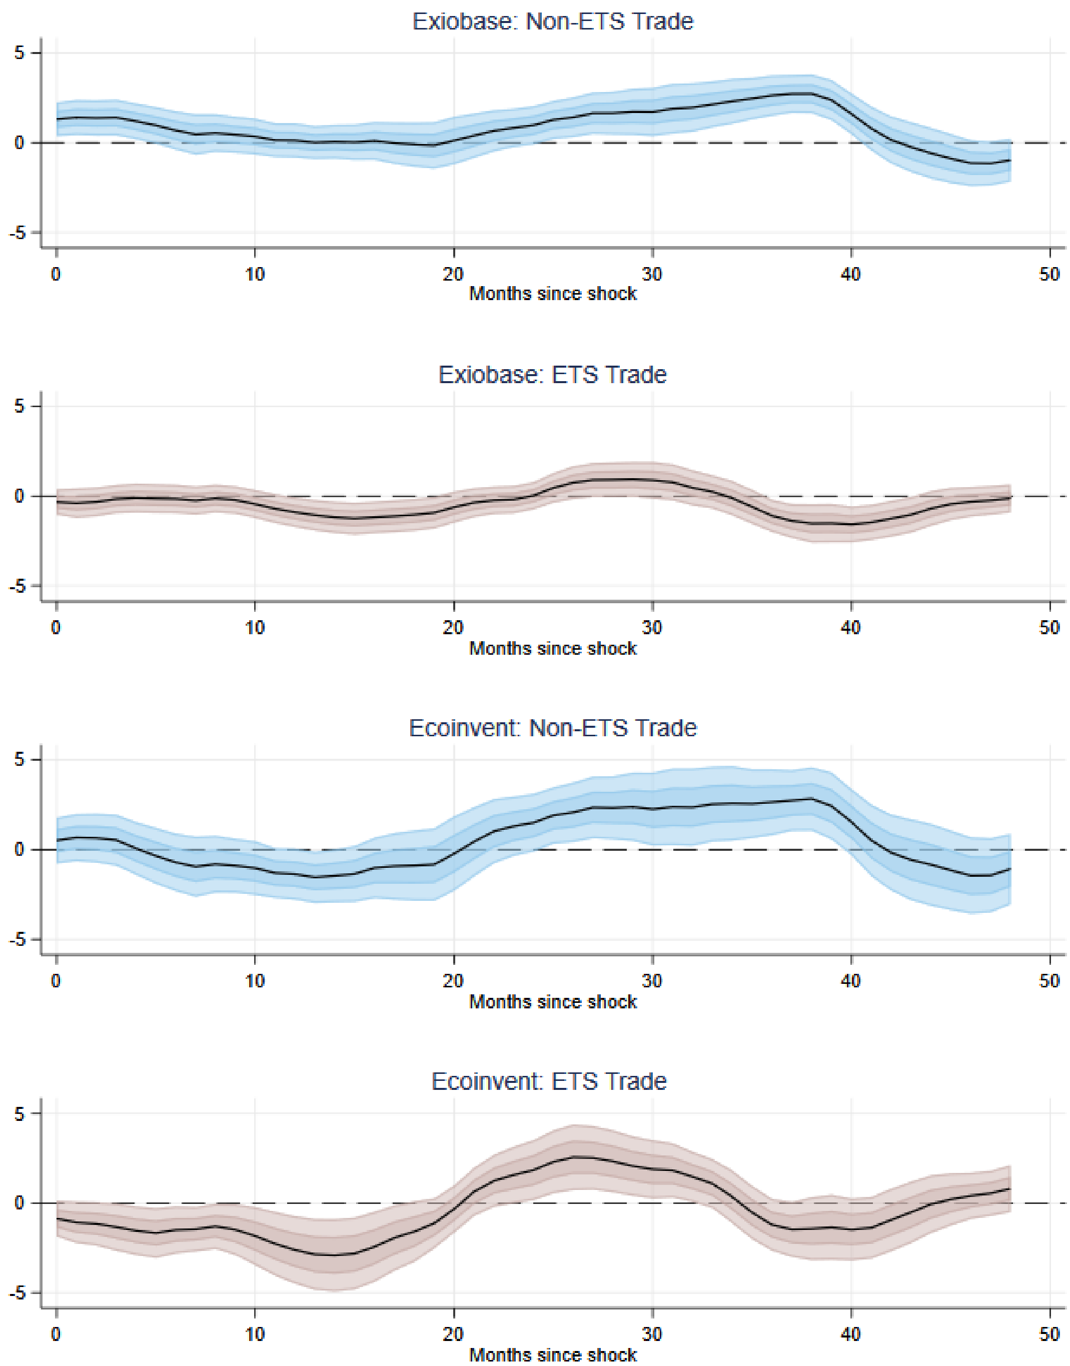

Notes: This figure displays the average firm response to a carbon price shock. The dependent variable is a 6-month moving average of the quantitiy of trade multiplied with average co2 content. In the upper two panels we use information on CO2 content from Exiobase while in the lower two panels we use information from Ecoinvent. We include firm-country, year and month fixed effects and control for 4 lags of the dependent variable. Standard errors are clustered at the firm and date level. The solid lines are the point estimates, and the shaded areas are the 95 and 68 percent confidence bands, respectively. Trade data underlying this graph comes from the customs data, emissions data comes from ecoinvent.

the products that do respond and that are not on the list are less carbon intense. We replicate the exemption regression using CO2 from Exiobase as the dependent variable in 5, confirming a similar pattern as before. The CO2 intensive goods on the exemption list also do not react to imported CO2 content.

6.4 Heterogeneity

To shed further light on potential mechanisms and as an additional robustness check, we divide the sample by 2-digit product codes and confirm considerable heterogeneity. A typical example of considerable carbon leakage in the literature, e.g., Aichele and Felbermayr (2015), is paper and paper products. We can confirm a more substantial and faster-than-average reply to this product category. A similar effect is true, for example, for Chemicals and Chemical products. Both product groups are relatively energy-intensive and, in their basic form, easy to trade and substitute, so a strong response is expected. In contrast, pharmaceutical products and pharmaceutical preparations show no or low signs of trade substitution. The high responsiveness of this product group would have been surprising as the sector overall features additional trade barriers because of regulation, international suppliers are usually limited, and production is not very carbon-intense. Overall, the results further support the notion that our results capture a response to a carbon price increase, and we do not capture other effects correlated with our carbon policy surprises that affect overall aggregate trading patterns. FIGURES PENDING

6.5 Discussion

Our results show that firms react to carbon price increases by substituting for imports from countries outside of the EU-ETS price domain. This channel is the intensive margin of firms' trading decisions. We do not assess if a firm is more likely to trade but analyze whether firms that already traded trade more. Furthermore, it captures the effect of a marginally changing CO_2 price already in place. It does not reflect the effect of the introduction of a CO_2 price. When a CO_2 price is implemented, firms form their expectation about long-term price developments, make long-term investment

Figure 5: Carbon leakage comparing exempt and non exempt products

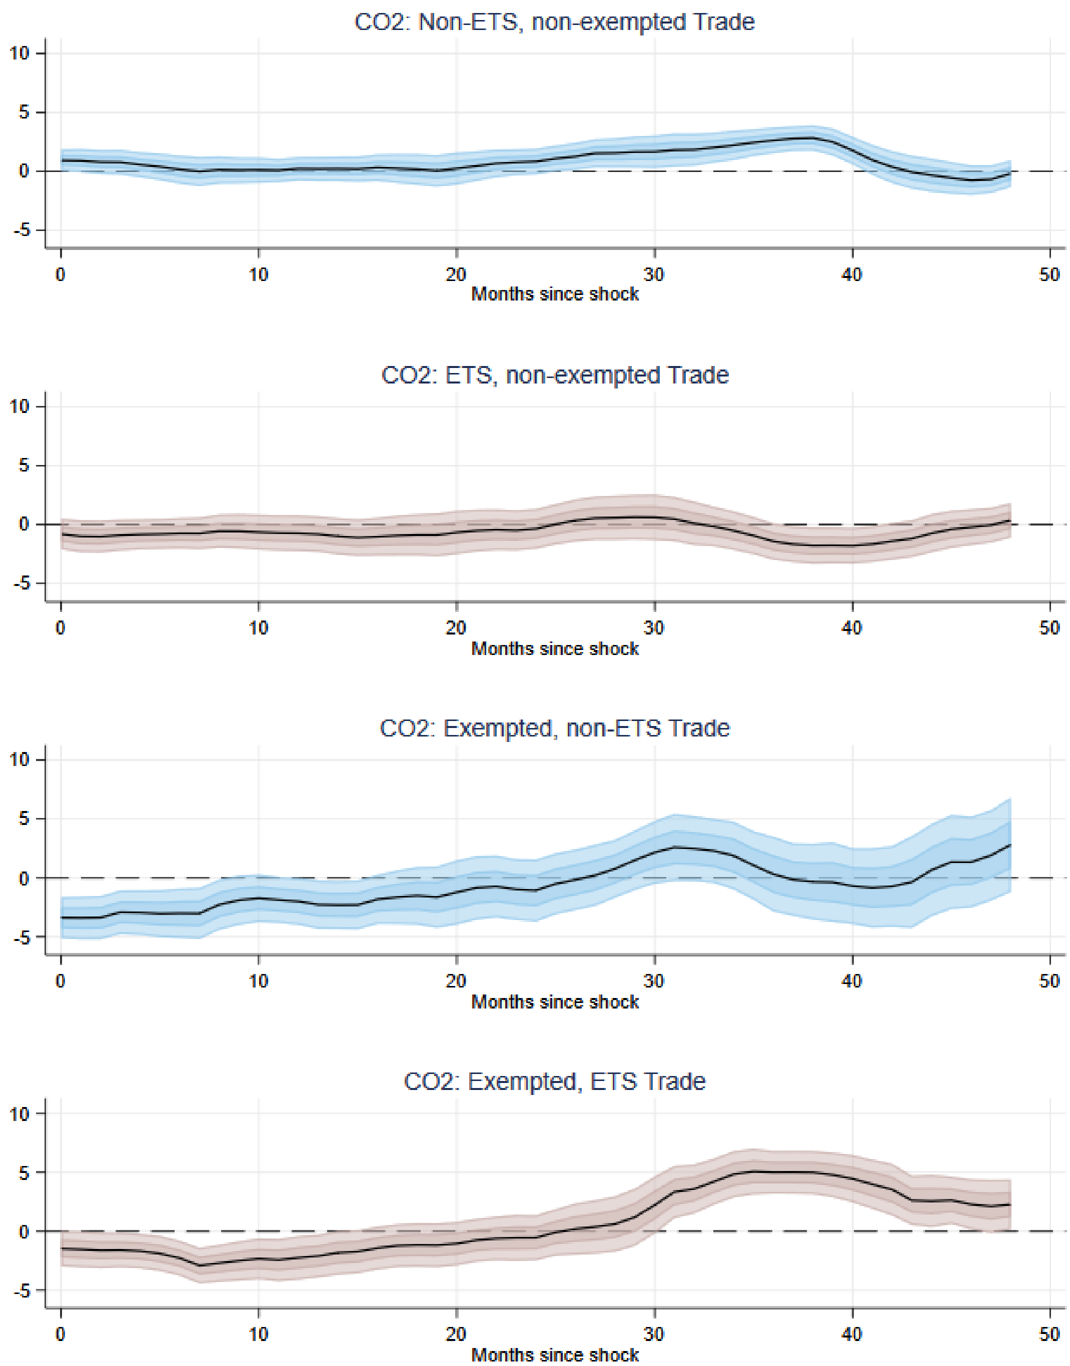

Notes: This figure displays the average firm response to a carbon price shock. he dependent variable is a 6-month moving average of the value of trade multiplied with average co2 content. We include firm-country, year and month fixed effects and control for 4 lags of the dependent variable. Standard errors are clustered at the firm and date level. The solid lines are the point estimates, and the shaded areas are the 95 and 68 percent confidence bands, respectively. Trade data underlying this graph comes from the customs data, emissions data comes from Exiobase 3

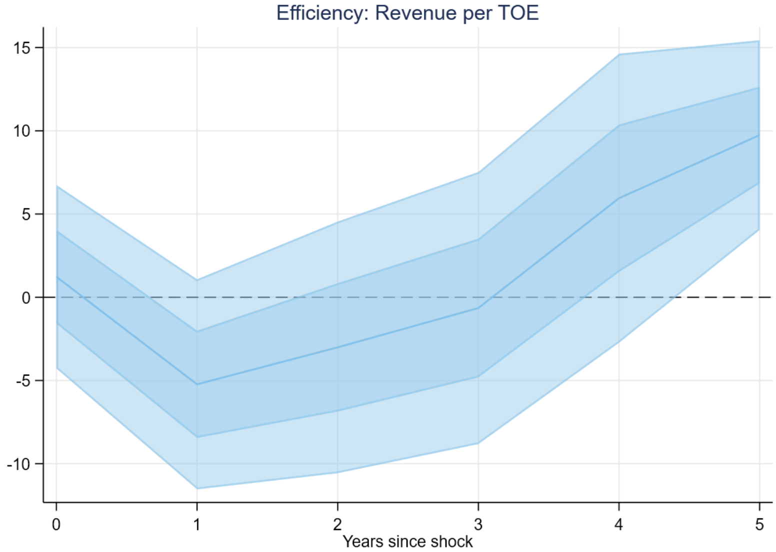

Figure 6: Revenue per tonne of oil equivalent (toe)

Notes: This figure displays the average firm response to a carbon price shock. The dependent variable is the share of revenue relative to energy used in toe. We include firm fixed effects and control for 4 lags of the dependent variable. Standard errors are clustered at the firm level. The solid lines are the point estimates, and the shaded areas are the 95 and 68 percent confidence bands, respectively. trade data underlying this graph comes from balance sheets (FICUS/FARE) and the energy consumption survey.

decisions abroad, and start looking for new suppliers abroad. These decisions may be implemented faster as a consequence of the carbon price changes we measure, but the price changes we measure are probably too small to trigger these decisions. One way to see our approach is to capture firms' reactions toward deviations from an anticipated price path and can, therefore, be seen as a lower bound for import substitution.

Furthermore, our coefficients capture the overall effect of trade after a marginal change in carbon prices. One might be inclined to think that a higher exposure of products to carbon prices should lead to a stronger response in trade. We tested several specifications using CO_2 intensity measures obtained both from Exiobase and Ecoinvent but could not confirm such a pattern. There are several potential explanations for this. An important factor is the general substitutability and overall trade elasticity of the product. It also matters how easy it is to generally source a specific product from abroad in a relatively short time. Another reason for the missing link between trade responsiveness and CO_2 intensity might lie in the EU's exemption and subsidy policies. As CO_2 and energy-intensive firms receive special treatment through free allocation and

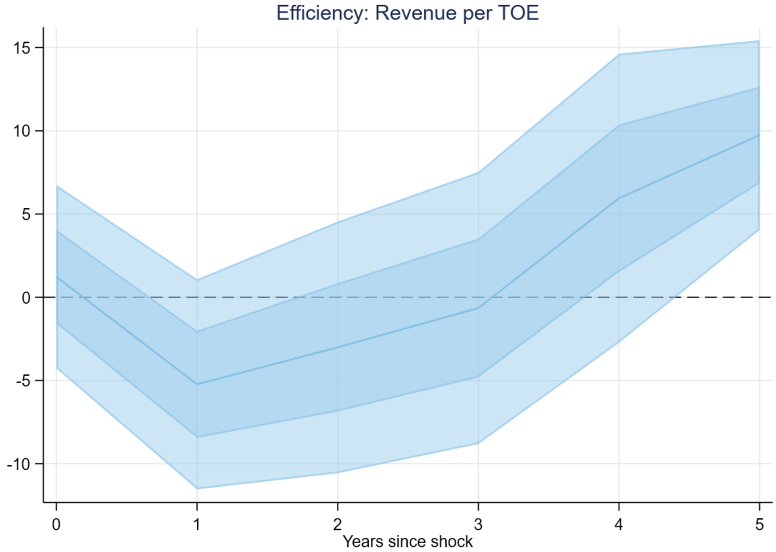

Figure 7: Revenue per tonne of oil equivalent (toe)

Notes: This figure displays the average firm response to a carbon price shock. The dependent variable is the share of revenue relative to energy used in toe. We include firm fixed effects and control for 4 lags of the dependent variable. Standard errors are clustered at the firm level. The solid lines are the point estimates, and the shaded areas are the 95 and 68 percent confidence bands, respectively. trade data underlying this graph comes from balance sheets (FICUS/FARE) and the energy consumption survey.

potentially direct energy subsidies, this mutes the potential channel of CO_2 intensity as our results show that, in particular, those products on the exemption list do not react.

Our results show a considerable time delay concerning the trade response and the effect peaks 30-40 months after the shock. There are several reasons why this is the case. First, timing frictions are naturally embedded in international trade, such as ordering, production, and shipping time. It also takes time for the supplier to expand their production or to find new suppliers abroad. Further, while the general CPI energy components react only in the short to medium term, the general shock propagation to firm prices takes longer as prices are sticky, and it takes longer to cascade through the supply chain. As shown in Hensel et al. (2023a), prices after a carbon price shock adjust for up to 13 quarters, which is consistent with the timeline of our trade estimates.

Nevertheless, how can one rationalize the missing long-term effect? There are three possible interpretations for the long-run effect of carbon price shocks: a more optimistic scenario, a more neutral assessment, and a more technical one. The optimistic one is that in reaction to a carbon price surprise, the economy starts to adapt, and within a five-year window after the shock, production technology becomes more efficient, leading to a phasing out of the shock in the long run. There is relatively robust evidence that innovation activity increases. Känzig (2022), for example, finds that patenting activity increases in reaction to a carbon shock. This seems to occur both on an energy producer side, as indicated by the relatively quick decline in the energy component of the CPI, but also on a firm level. In Figure 7, we have shown suggestive evidence that revenue per energy input increases in reaction to a shock. This cannot be due to import substitution as the effect has already faded out when we detect the effect. For the neutral explanation, we again think in a framework that we caption the firm's reaction towards deviations from an anticipated price path. In the long run, prices come closer to the original projection that anticipated an increasing carbon price, and it just came earlier than expected. Consequently, companies increased the transition speed, which is reflected in the medium-term response. The more technical view is that these surprise shocks tend to phase out over time. This can for example be seen in Figure 1 and is known to happen for the more frequently used monetary policy surprises. Whether this happens exactly because the economy adjusts to the new environment or whether this is a technical feature of the shocks is hard to say.

7 Exports

RESULTS PENDING

8 Conclusion

Climate change is one of the major challenges of our time, and we are already late in applying the necessary measures. The only question is now, how much global warming will there be? Every CO_2 emission we can reduce will reduce the costs associated with climate change. One of the most efficient ways is to impose a price on carbon, for example, through an emissions trading system like the EU-ETS. However, the price of carbon is only efficient if it cannot be avoided by buying products from abroad. In this project, we have investigated empirically the extent to which firms' trading patterns change in reaction to changes in carbon prices.

To isolate exogenous variation in carbon prices, we rely on the shocks identified by Känzig (2022). He identified significant institutional events that influenced the supply of emission allowances. He then isolates the surprise by looking at carbon future price changes in a tight window around the event. Since the surprises are only a partial measure, they are further improved by using them as an input in a proxy-VAR. We use the resulting shocks as an independent variable in our local projections à la Jordà (2005). This technique allows us to trace the firms' responses to a carbon price shock over time.

We use French firm-level trade data to show that in reaction to a carbon price shock trade with non-European countries increases but trade with other European countries is mostly unaffected. Trade with non-European countries increases slightly after a shock but then remains mostly unaffected for the first one and a half years. For the next two years, trade with non-European countries is elevated. A carbon shock that increases energy prices by 1% increases trade with non-European countries by up to 4% at a given horizon. However, after five years trade volumes have converged back to pre-shock levels. We then use two different sources to attribute CO_2 emissions to these trade flows, Exiobase and Ecoinvent. We find a very similar pattern for CO_2 emissions embedded in these trade flows, although, lower in magnitude. We further provide suggestive evidence that Firms adjust to the shock by increasing their revenue per energy unit used.

The European Union defined certain industries and products as being at significant risk of carbon leakage and implemented measures to protect them. We analyze how the reaction of these industries/products differs compared to the rest of the economy. We find that the response to a carbon shock differs significantly. Imports from non-ETS countries do not increase for products that were identified as being at risk of carbon leakage. However, imports of these products from other European countries do indeed increase. We conclude that this is suggestive evidence that the measure actually helped attenuate carbon leakage by leading to a reallocation process within the European Union.

We further show that the trade reaction to carbon shocks differs by industry. We show that certain energy-intensive industries such as for example paper and pulp production react strongly to the shock while others like pharmaceutical products and pharmaceutical preparations show no reaction to a carbon shock. We take this as evidence for the validity of the shock.

In the past, several studies on the EU-ETS found no or only a little carbon leakage. We propose a possible explanation, namely that carbon leakage is only temporary and fades out in the long run. Further, special policies targeting products and sectors at risk of carbon leakage seem to be effective at preventing import substitution and consequently carbon leakage. This suggests that the implementation of an Emissions Trading System has to be done carefully with the threat of carbon leakage in mind. However, we believe that our estimate should be treated as a lower bound as we do not capture the effects of introducing a Carbon Price in the first place but rather of marginally changing a price that is already in place. While one can argue that this is the relevant policy measure at hand right now more research is needed to capture the overall effect. However, independent of this re results show the potential positive effects of a carbon border adjustment which could not only address the effects of temporary carbon leakage measured in this paper but would at the same time generally level the playing field between countries with a carbon price and without.

References

Aichele, R. and Felbermayr, G. (2015). Kyoto and Carbon Leakage: An Empirical Analysis of the Carbon Content of Bilateral Trade. *The Review of Economics and Statistics*, 97(1):104–115.

Bergounhon, F., Lenoir, C., and Mejean, I. (2018). A guideline to French firm-level trade data. Technical report.

Berthold, B., Cesa-Bianchi, A., Pace, F. D., and Haberis, A. (2023). The Heterogeneous Effects of Carbon Pricing: Macro and Micro Evidence. Technical report.

Burstein, A., Carvalho, V., and Grassi, B. (2020). Bottom-up Markup Fluctuations. Technical report, National Bureau of Economic Research, Cambridge, MA.

De Ridder, M., Grassi, B., and Morzenti, G. (2022). The Hitchhiker's Guide to Markup Estimation *. Technical report.

Dechezleprêtre, A., Gennaioli, C., Martin, R., Muûls, M., and Stoerk, T. (2022). Searching for carbon leaks in multinational companies. *Journal of Environmental Economics and Management*, 112.

di Giovanni, J., Levchenko, A. A., and Mejean, I. (2014). Firms, Destinations, and Aggregate Fluctuations. *Econometrica*, 82(4):1303–1340.

Hensel, J., Mangiante, G., and Moretti, L. (2023a). Carbon Pricing and Inflation Expectations: Evidence from France. Technical report.

Hensel, J., Mangiante, G., and Moretti, L. (2023b). Carbon Pricing and Inflation Expectations: Evidence from France. *SSRN Electronic Journal*.

Jordà, (2005). Estimation and Inference of Impulse Responses by Local Projections. *American Economic Review*, 95(1):161–182.

Känzig, D. R. (2022). The unequal economic consequences of carbon pricing. Technical report.

Mangiante, G. (2023). The Geographic Effects of Carbon Pricing. Technical report.

Montiel Olea, J. L., Stock, J. H., and Watson, M. W. (2018). Inference in Structural Vector Autoregressions Identified With an External Instrument. Technical report.

Muûls, M., Martin, R., Colmer, J., and Wagner, U. (2022). Does Pricing Carbon Mitigate Climate Change? Firm-Level Evidence from the European Union Emissions Trading Scheme. Technical report.

Nakamura, E. and Steinsson, J. (2018). High-Frequency Identification of Monetary Non-Neutrality: The Information Effect*. *The Quarterly Journal of Economics*, 133(3):1283–1330.

Newey, W. K. and West, K. D. (1987). Hypothesis Testing with Efficient Method of Moments Estimation. *International Economic Review*, 28(3):777.

Stadler, K., Wood, R., Bulavskaya, T., Södersten, C. J., Simas, M., Schmidt, S., Usubiaga, A., Acosta-Fernández, J., Kuenen, J., Bruckner, M., Giljum, S., Lutter, S., Merciai, S., Schmidt, J. H., Theurl, M. C., Plutzar, C., Kastner, T., Eisenmenger, N., Erb, K. H., de Koning, A., and Tukker, A. (2018). EXIOBASE 3: Developing a Time Series of Detailed Environmentally Extended Multi-Regional Input-Output Tables. *Journal of Industrial Ecology*, 22(3):502–515.

Stock, J. H. and Watson, M. W. (2018). Identification and Estimation of Dynamic Causal Effects in Macroeconomics Using External Instruments. In *Economic Journal*, volume 128, pages 917–948. Blackwell Publishing Ltd.

Venmans, F., Ellis, J., and Nachtigall, D. (2020). Carbon pricing and competitiveness: are they at odds? *https://doi.org/10.1080/14693062.2020.1805291*, 20(9):1070–1091.

Yu, B., Zhao, Q., and Wei, Y. M. (2021). Review of carbon leakage under regionally differentiated climate policies. *Science of The Total Environment*, 782:146765.

A Additional information

A.1 Construction of shocks

The measurement of the surprise series, although very clever, is imperfect. There will be some measurement error, for example, because not all relevant events have been identified or the window around the announcement has to be relatively wide (1 Day Känzig (2022)). Hence, following Stock and Watson (2018) the surprise series is used as an external instrument in structural VAR. In this section, we will briefly outline the approach that was performed by Känzig (2022) to arrive at the shock series, but is central for the understanding of our paper.

We use the standard notation in the VAR literature. As a quick reminder: vectors and matrices are denoted in bold e.g. $\boldsymbol{y}_{t-p}$ where the subscript t denotes the year and p denotes the lag of a variable. Individual elements are written as $y_{i,t}$, where i denotes the row. $\boldsymbol{y}_t$ denotes the $n \times 1$ vector of endogenous variables and $\boldsymbol{B}_1$ to $\boldsymbol{B}_p$ are coefficient matrices of size $n \times n$. $\boldsymbol{u}_t$ is the $n \times 1$ vector of reduced form innovations and $Var(\boldsymbol{u}_t) = \Sigma$.

A standard VAR model takes the following form

$$\boldsymbol{y}_t = \boldsymbol{B}_0 + \boldsymbol{B}_1 \boldsymbol{y}_{t-1} + ... + \boldsymbol{B}_p + \boldsymbol{y}_{t-p} + \boldsymbol{u}_t \tag{4}$$

Under the assumption that the VAR is invertible, $\boldsymbol{u}_t$ can be defined as

$$\boldsymbol{u}_t = \boldsymbol{S}\boldsymbol{\epsilon}_t \tag{5}$$

Where $\boldsymbol{S}$ is an $n \times n$ non-singular matrix. The structural shocks $\boldsymbol{\epsilon}_t$ are assumed to be serially and mutually uncorrelated. Without loss of generality we order our shock of interest first in $\boldsymbol{\epsilon}_t$, thus it is denoted as ϵ_{1_t}. The scales of $\boldsymbol{\epsilon}_{1,t}$ and $\boldsymbol{S}$ are not separately identified. It thus makes sense to scale such that they have a meaningful economic interpretation. In our case, they are scaled such that a one unit increase in ϵ_{1_t} increases household energy prices by 1%.

To be used as an external instrument z_t, our surprise series is required to be relevant, thus being correlated with the underlying structural shock of interest and exogenous, in other words, uncorrelated with other shocks in the model Montiel Olea et al. (2018). Thes conditions can be written as

$$\mathbb{E}[z_t \epsilon_{1,t}] = \alpha \neq 0 \tag{6}$$

$$\mathbb{E}[z_t \epsilon_{j,t}] = 0 \text{ for } j \neq 1 \tag{7}$$

Using these assumptions the instrument can be used to recover the structural shock of interest $\epsilon_{1,t}$ [5].

B Additional figures and tables

[5]See Stock and Watson (2018) page 4 for the derrivation